

Notes on Transformer Models

An exposition for test fixtures.

Abstract. This short note summarises several widely cited results in modern machine learning. It exists only as a test fixture for the citation-cleaner tool.

1. Introduction. The Transformer architecture (Vaswani et al., 2017) revolutionised sequence modelling. Earlier work established the foundations of deep learning (LeCun, Bengio, and Hinton, 2015) and ImageNet-scale image classification (Krizhevsky, Sutskever, and Hinton, 2012).

2. Discussion. Later work on BERT (Devlin et al., 2019) showed that Transformer encoders transfer well to many NLP tasks.

References

- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems* 30.
- LeCun, Y., Bengio, Y., and Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436-444.
- Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. *Advances in Neural Information Processing Systems* 25.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL-HLT 2019*.
- He, K., Zhang, X., Ren, S., and Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Smith, J. and Jones, K. (2024). A unified theory of toaster cognition. *Journal of Imaginary Sciences*, 42(7), 999-1013.
- Doe, J., Roe, R., and Moe, M. (2023). Quantum entanglement of breakfast cereals. *Proceedings of the International Conference on Nonexistent Results*, 17-29.